Stanford
ONLINE

DeepLearning.AI

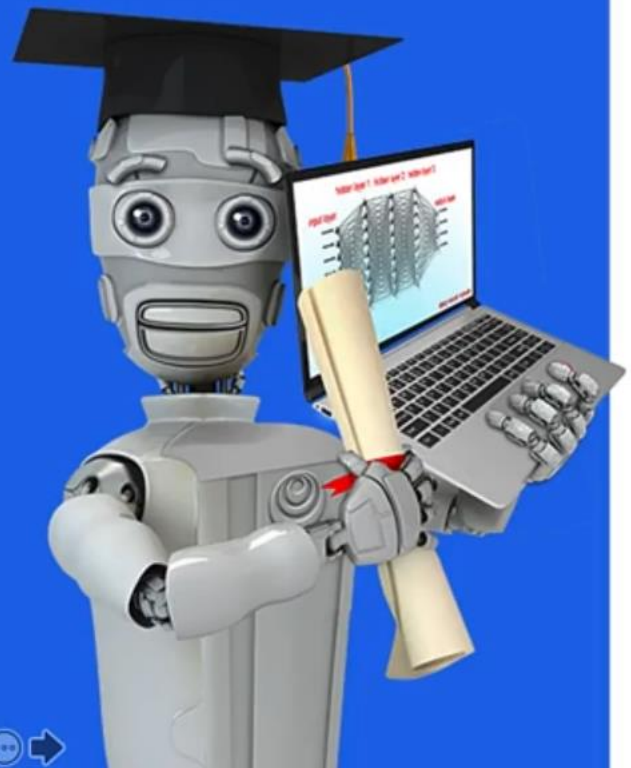

Linear Regression with One Variable

Cost Function

Training set

features | targets

size in feet² (x)	price \$1000's (y)
2104	460
1416	232
1534	315
852	178
...	...

Model: $f_{w,b}(x) = wx + b$

w, b: parameters

coefficients

weights

What do w, b do?

$$f_{w,b}(x) = wx + b$$

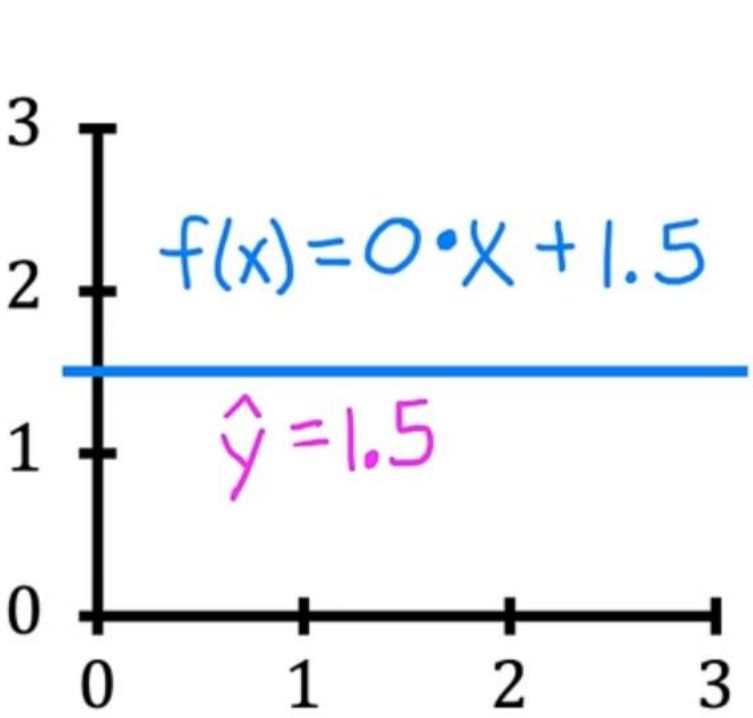

→ $w = 0$

→ $b = 1.5$

(y-intercept)

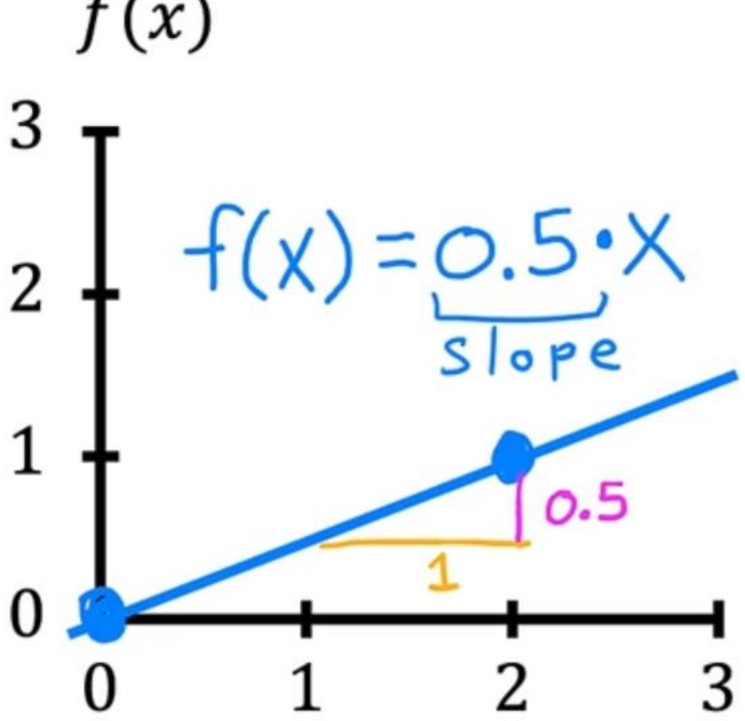

→ $w = 0.5$

→ $b = 0$

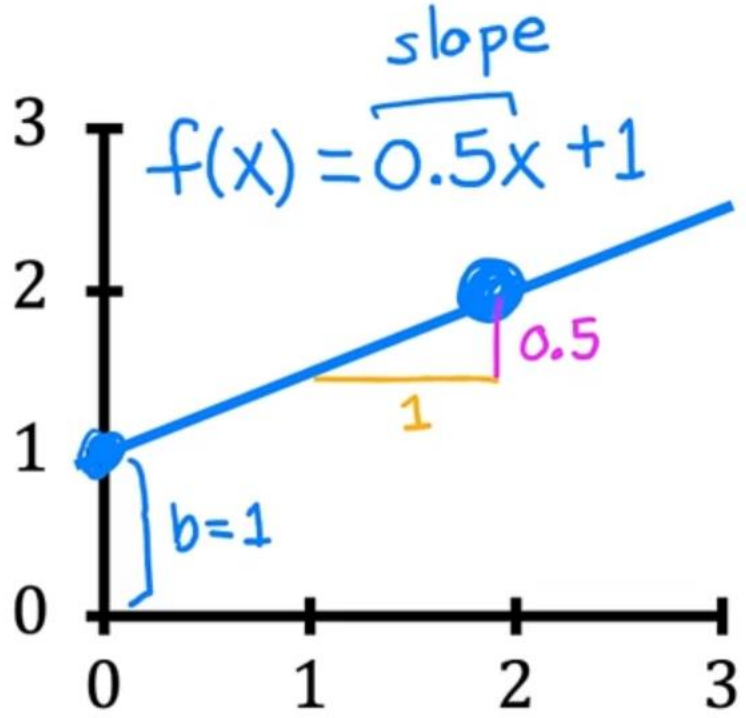

→ $w = 0.5$

→ $b = 1$

Cost function: Squared error cost function

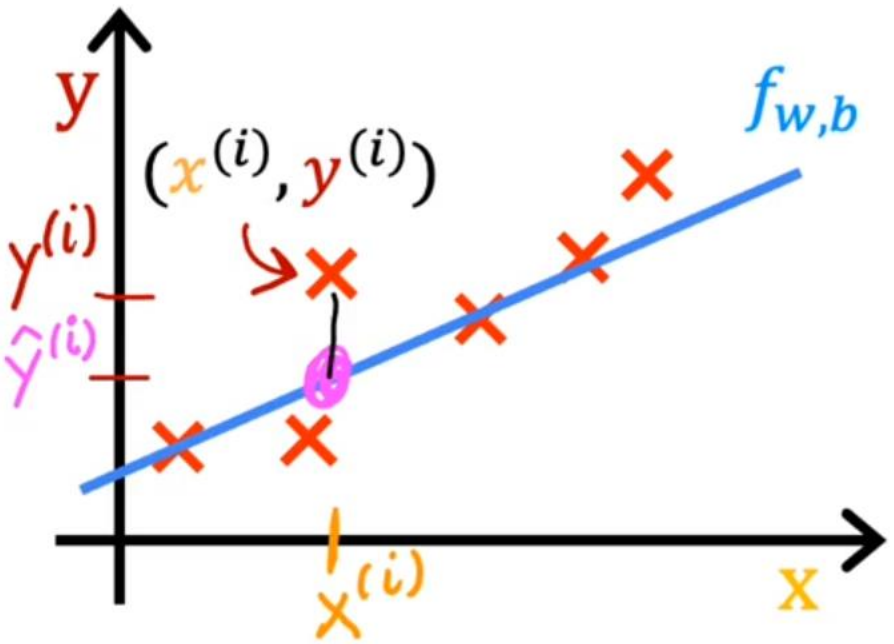

$$\hat{y}^{(i)} = f_{w,b}(x^{(i)})$$

$$f_{w,b}(x^{(i)}) = wx^{(i)} + b$$

$$J(w,b) = \frac{1}{2m}\sum_{i=1}^{m}\left(\hat{y}^{(i)} - y^{(i)}\right)^2$$

error

m = number of training examples

$$J(w,b) = \frac{1}{2m}\sum_{i=1}^{m}\left(f_{w,b}(x^{(i)}) - y^{(i)}\right)^2$$

intuition (next!)

Find w, b:

$\hat{y}^{(i)}$ is close to $y^{(i)}$ for all $(x^{(i)}, y^{(i)})$.