Stanford
ONLINE

DeepLearning.AI

Training Linear Regression

Gradient Descent

Have some function $J(w,b)$ for linear regression or any function

Want $\min_{w,b} J(w,b)$ $\min_{w_1,\ldots,w_n,b} J(w_1,w_2,\ldots,w_n,b)$

Outline:

- Start with some w,b (set $w=0$, $b=0$)
- Keep changing w,b to reduce $J(w,b)$
- Until we settle at or near a minimum

J not always

may have >1 minimum

gradient descent

$J(w,b)$

not squared error cost
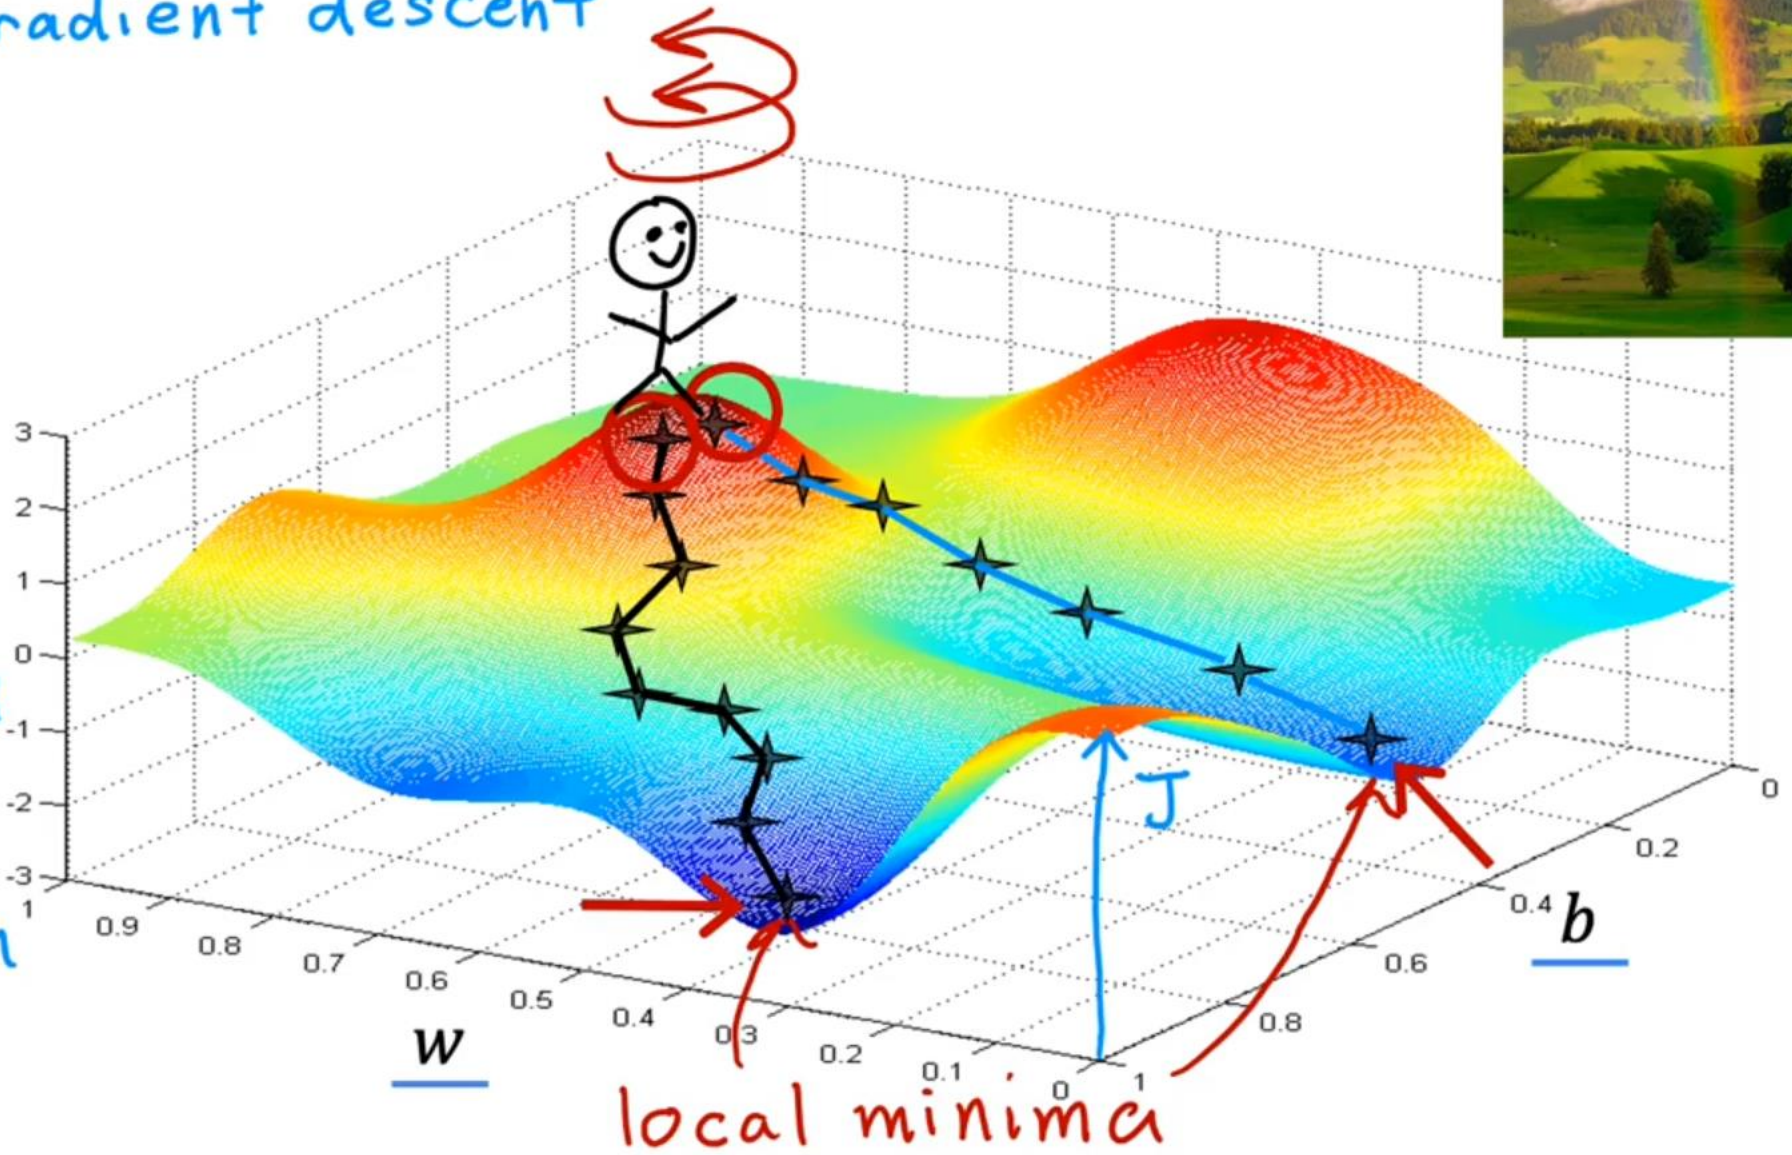
not linear regression